


pyNTM

Training Module 1 - intro to pyNTM and modeling basics

python3_Network_Traffic_Modeler (pyNTM)

Tim Fiola




Versioning

- ▶ Training v1.0
- ▶ pyNTM v1.5

Agenda

- ▶ Problem statement
- ▶ Network modeling provides value
- ▶ Simulations
- ▶ Network modeling use cases
- ▶ We need open source tools in this space
- ▶ What is pyNTM?
- ▶ pyNTM and *networkx*
- ▶ Why is pyNTM helpful?
- ▶ pyNTM mechanics
- ▶ pyNTM features
- ▶ Wide area network modeling options
- ▶ Who is pyNTM for?



Problem Statement - Understanding the wide area network during failure states and maintenances and how to grow the network

- ▶ In a large, meshy network, it becomes difficult to understand how a given failure will truly impact interface utilization in other parts of the network
 - ▶ leads to educated guessing and general *rules of thumb*
- ▶ WAN capacity cannot be solved simply throwing money at the problem
 - ▶ WAN circuits are expensive
 - ▶ WAN circuits are not always available
- ▶ Capital in the WAN must be efficiently allocated

Network modeling provides strategic value

- ▶ Modeling allows unique, data-based understanding of how network will behave during
 - ▶ Failover
 - ▶ Changes in the traffic matrix
 - ▶ Changes in topology, such as adding RSVP mesh or changing a link metric
- ▶ This understanding prevents
 - ▶ overbuilding WAN links, which strands capital
 - ▶ Underbuilding WAN links, which increases risk

A network model provides value for people in the following roles (to name a few)

- ▶ Capacity Planner
 - ▶ Plan network to optimize latency, cost, simplest topology, etc
- ▶ Network Engineer
 - ▶ Test different topologies
- ▶ Anyone working a maintenance
 - ▶ Simulate the effects of taking down a router (Node) for a maintenance
- ▶ Anyone with interest in network performance

A network model has two inputs

► Traffic Matrix

- Each entry describes a *demand*
- Each demand has
 - *magnitude*, which describes how much traffic is in that demand
 - A source and destination node
- The traffic matrix for a network will vary throughout the day, month, season, etc
- Getting good traffic matrices can be challenging
 - Understanding your network's traffic matrices allows for truly effective engineering and planning

► Topology

- Layer 3 nodes
- Circuits (comprised of 2 interfaces) between layer 3 nodes
- Shared Risk Link Groups (SRLGs)
- RSVP LSPs

Sample traffic matrix (Mbps)

Destination \ Source	A	B	C
A	-	45	120
B	60	-	
C	75	150	-

This example traffic matrix shows traffic sourced from Node A destined to Node C with a magnitude of 120Mbps

How will this traffic transit the network?

Network modeling provides *simulation* capability

- ▶ Applying the traffic matrix to the topology and converging the model produces a *simulation*
 - ▶ The *modeling engine* governs behaviors of demands, LSPs, etc in the model
 - ▶ pyNTM has a modeling engine
- ▶ For a given day, you can produce a simulation for different parts of the day by creating a traffic matrix that reflects source and destination traffic entries for each part of the day
 - ▶ What happens during a given failure if it were to occur at different parts of the day?
 - ▶ What is the best time to conduct a maintenance on a given router?
 - ▶ Where is the best place to augment the network to best handle our holiday traffic matrices?

Wide area network modeling vs legacy techniques

Non-modeling techniques

- ▶ Rules of thumb, such as augmenting a circuit when it reaches 50% utilization, don't guarantee that a given circuit will be able to handle the events *you are interested in*
- ▶ Planning an auto-bandwidth RSVP network adds additional complexity
 - ▶ The demands a link handles can change throughout the day/week/season

Modeling

- ▶ Simulations allow for
 - ▶ Understanding how your network traffic will behave during a failure event
 - ▶ Understanding how your network will behave with additional traffic
 - ▶ Better understanding of how RSVP LSP meshes will behave
- ▶ Understanding your network allows you to
 - ▶ Efficiently allocate capacity/capital where it's really needed
 - ▶ Plan for and understand events you care about
- ▶ At a minimum, allows you to make a more educated estimation

Some example use cases for network model



Understanding
current network
topology

How many ECMP paths
does a given demand
take across the network?



Understanding
failover by modeling
failures

Link(s)
Node(s)
Shared risk link
group(s) (SRLG)



Understanding where
it makes sense to
augment a network


Deploy capital
where it's most
needed
Don't strand capital

More example use cases for network model



Understanding how
changes in the
network affect
traffic flow

More/less traffic
Adding capacity to
existing link(s)
New link(s)
Metric changes



RSVP
Implementations

Adding RSVP
overlay to IGP
network
Adding/removing
parallel LSPs
Failover

The need

- ▶ Open-source tools that allow programmatic network modeling and simulation
- ▶ Specifically, there are two needed components
 - ▶ Open source modeling engines
 - ▶ Open source tools to create reasonable traffic matrices
- ▶ Nice to have: open source GUI for visualization

So, what is pyNTM?

- ▶ pyNTM is the python3 Network Traffic Modeler
- ▶ pyNTM is an **open source** WAN *modeling engine*
- ▶ Applies a traffic matrix to a network topology to route traffic as the network would
 - ▶ Uses *networkx* module to get the topology path info
 - ▶ *networkx* is a GREAT tool to get path info in a topology . . .
 - ▶ . . . but there's more to modeling than just path info

Networkx and pyNTM roles in pyNTM simulations

Demand routes (**networkx**)



RSVP routes (**networkx**)




Interface utilization (**pyNTM**)

Interface reserved bandwidth (**pyNTM**)


Demand path's interfaces (**pyNTM**)

RSVP LSP path's interfaces (**pyNTM**)



Networkx and pyNTM with ECMP traffic

- ▶ Networkx is great to find all the unique paths through a topology
- ▶ However, you can't just model a demand with ECMP by splitting traffic evenly across all the unique paths
 - ▶ That would be *end-to-end* load balancing
- ▶ In the example to the right, a demand with a magnitude of 12 enters the topology
- ▶ Spreading the traffic evenly across the 3 unique paths results in the traffic spread shown



Unique paths are:


A-B-D-G

A-C-E-G


A-C-F-G

Networkx and pyNTM with ECMP traffic

- ▶ pyNTM models *hop-by-hop* ECMP across the 3 unique paths
 - ▶ This is how OSPF and ISIS load balance
- ▶ In the example shown, a demand with a magnitude of 12 enters the topology
- ▶ Hop-by-hop load balancing results in the traffic spread shown
- ▶ This traffic spread across the 3 paths is very different than the end-to-end load balancing traffic spread




Why is pyNTM helpful?



- ▶ PyNTM leverages path information from *networkx* in a *network state-specific context*, allowing for modeling of *network-specific* state:
 - ▶ Modeling utilization on interfaces
 - ▶ Modeling traffic demands consuming interface bandwidth
 - ▶ Modeling RSVP LSPs consuming reservable interface bandwidth
 - ▶ Associating a traffic demand with the specific interfaces along the path(s)
 - ▶ Determining the available path(s) that have a given amount of reservable bandwidth
- ▶ pyNTM APIs allow for programmatic network modeling capability

Why is pyNTM helpful? (continued)



- ▶ pyNTM allows users to easily modify the network topology and determine alternate network state based on that change
 - ▶ Failing layer3 Nodes, Circuits, SRLGs, etc
 - ▶ Adding new Nodes, Interfaces, traffic Demands, etc to the topology
 - ▶ Adding new/additional auto-bandwidth LSPs to the network
- ▶ pyNTM is specifically designed to easily relate the following items:
 - ▶ Traffic Demands
 - ▶ RSVP LSPs
 - ▶ Path info
 - ▶ Interfaces
 - ▶ Nodes

pyNTM Mechanics

- ▶ *pyNTM Model* object describes the network topology
- ▶ The Model object contains the following objects:
 - ▶ Interfaces/Circuits
 - ▶ Nodes
 - ▶ Demands (traffic)
 - ▶ RSVP LSPs
 - ▶ Shared Risk Link Groups
- ▶ The Model object applies the traffic matrix to the topology, allowing the traffic to flow as it would in a real network
 - ▶ This produces a simulation
- ▶ Valuable data can be mined from simulation results


pyNTM features (as of v1.5)

- ▶ IGP (OSPF/ISIS)
- ▶ RSVP LSPs
 - ▶ Supports full mesh
 - ▶ IGP shortcuts is on the roadmap
 - ▶ Auto-bandwidth
 - ▶ Fixed/manually-assigned setup bandwidth
- ▶ Shared Risk Link Groups (SRLGs)
- ▶ Feature requests are accepted on GitHub!
- ▶ Currently only supports modeling a single link between layer 3 nodes
 - ▶ Modeling multiple/parallel links between nodes may incur a large performance cost
 - ▶ We have top people looking into that problem . . . *TOP . . . PEOPLE*

Options for modeling/simulation

Feature	Commercial Options	pyNTM
Cost	\$tens-hundreds of thousands	\$0
APIs for programmatic modeling	Y	Y
Includes capability to create traffic matrix	Y	N
Sophisticated GUI for visualization	Y	N
Open Source	N	Y
Dependent on vendor	Y	N

So who is pyNTM for?



- ▶ If your org/company can generate a reasonable traffic matrix
 - ▶ Access to data scientists
 - ▶ PMACCT and NetFlow
 - ▶ Forecasted traffic demands
- ▶ If your org has basic python coding skills
- ▶ If your org does not want to rely on and/or manage external modeling vendors

pyNTM provides the open source modeling and simulation engine and can help you today

Next steps

- ▶ Continue with the follow-on pyNTM training modules that explain
 - ▶ How to get pyNTM from github or pypi
 - ▶ How to set up pyNTM in a virtual environment
 - ▶ How to use the example scripts included in the github repository
 - ▶ How to use pyNTM to solve for some basic, but very powerful, use cases

FIN

Backup slides

The value proposition for network modeling

- ▶ In a large, meshy network it becomes very difficult to understand what traffic flow will look like for a given failure event
 - ▶ This difficulty causes operators to augment and/or turn up additional, expensive WAN circuits in response, using general rules of thumb which don't guarantee that a given circuit will be able to handle failure events that you are interested in
 - ▶ augmenting a circuit when it reaches 50% utilization
 - ▶ This difficulty causes network planners to spend hours/days/months understanding how their network should be planned to deal with any given event (new traffic, failures, etc.)
- ▶ WAN capacity cannot be solved by throwing money at the problem
 - ▶ WAN circuits are expensive
 - ▶ WAN circuits are not always available
- ▶ Simulations allow for
 - ▶ Understanding how your network traffic will behave during a failure event
 - ▶ Understanding how your network will behave with additional traffic

The value proposition for network modeling (continued)

- ▶ Understanding your network allows you to
 - ▶ Smartly augment your network
 - ▶ Don't augment where it's not needed (stranding capital) for the failure scenarios you care about
 - ▶ Ensure you augment where it will be needed for the failure scenarios you care about
 - ▶ Understand how new traffic will affect the existing network in steady state and failover

Objects you'll find in a network model

- ▶ Interface
 - ▶ Represents an interface on one end of a Circuit
- ▶ Circuit
 - ▶ Two-way link between 2 Nodes
 - ▶ Comprised of an Interface on each side
 - ▶ Each component Interface sends traffic in a single direction (toward remote Node)
- ▶ Node
 - ▶ Layer 3 device
- ▶ Demand
 - ▶ Represents traffic load from a source Node to a destination Node
 - ▶ Has a magnitude that represents the amount of traffic
- ▶ RSVP LSP
- ▶ Shared Risk Link Group (SRLG)
 - ▶ A grouping of Interfaces and/or Nodes that share a similar risk
 - ▶ When the SRLG fails, all component objects fail